Tech United Eindhoven @Home
2017 Team Description Paper

M.F.B. van der Burgh, J.J.M. Lunenburg, R.P.W. Appeldoorn, R.W.J. Wijnands,
T.T.G. Clephas, M.J.J. Baeten, L.L.A.M. van Beek, R.A. Ottervanger, H.W.A.M. van Rooy
and M.J.G. van de Molengraft

Eindhoven University of Technology,
Den Dolech 2, P.O. Box 513, 5600 MB Eindhoven, The Netherlands
`http://www.techunited.nl`, `techunited@tue.nl`, `https://github.com/tue-robotics`

Abstract. This paper provides an overview of the main developments of the Tech United Eindhoven RoboCup@Home team. Tech United uses an advanced world modeling representation system called the Environment Descriptor that allows straight forward implementation of localization, navigation, exploration, object detection & recognition, object manipulation and robot-robot cooperation skills. Recently developments are improved object detection via deep learning methods, a generic GUI for different user levels, improved speech recognition and improved natural language interpretation. These developments are done on AMIGO and SERGIO and will now also be implemented on the Toyota HSR.

1 Introduction

Tech United Eindhoven[1] is the RoboCup student team of Eindhoven University of Technology[2] that (since 2005) successfully competes in the robot soccer Middle Size League (MSL) and later (2011) joined the ambitious @Home League. The Tech United @Home team is the vice champion of RoboCup 2016 in Leipzig and the reigning European Champion of the 2016 RoboCup European Open. The robot soccer middle-size Tech United team has an even greater track record with 3 world championship titles. See the Tech United website for more results. Tech United Eindhoven consists of (former) PhD and MSc students and staff members from different departments within the Eindhoven University of Technology.

This Team Description Paper is part of the qualification package for RoboCup 2017 in Nagoya, Japan and describes the current status of the @Home activities of Tech United Eindhoven. The main achievement of our long-term development is our generic world model ED. Recent developments are improved object detection via deep learning methods, a generic GUI for different user levels, improved speech recognition and improved natural language interpretation.

2 Environment Descriptor (ED)

The TUe Environment Descriptor (ED) is a Robot Operating System (ROS) based 3D geometric, object-based world representation system for robots. In itself ED is database system that structures multi-modal sensor information and represents this in an object-based world representation that can be utilized for robot localisation, navigation, manipulation and interaction functions. See Figure 1 for a schematic overview of ED. ED is used on our robots AMIGO and SERGIO

---

[1] `http://www.techunited.nl`

[2] `http://www.tue.nl`

in the open @Home league and will be used on the Toyota HSR in the DSPL. In previous years, developments have been focussed towards making ED platform independent. As a results ED had been used on the PR2 system, Turtlebot and Dr. Robot systems (X80). ED is one re-usable
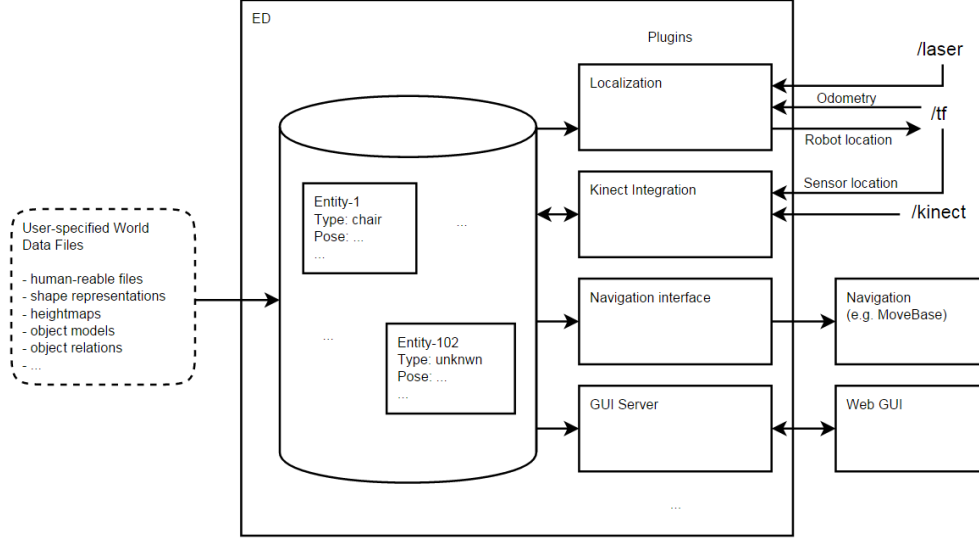

Fig. 1: schematic overview of TUe Environment Descriptor.

environment description that can be used for a multitude of needed functionalities. Instead of having different environment representations for localization Adaptive Monte Carlo Localization (AMCL), navigation (MoveBase), manipulation (MoveIt!), interaction, etc.. An improvement in this single, central world model will reflect in the performances of the separate robot capabilities. It omits updating and synchronization of multiple world models. At the moment different ED modules exist that enable robots to localize themselves, update positions of known objects based on recent sensor data, segment and store newly encountered objects and visualize all this through a web-based GUI, illustrated in Figure 9.

2.1 Localization, Navigation and Exploration

The ED-localization[3] plugin implements AMCL based on a 2D render from the central world model. In order to navigate, a model of the environment is required. This model is stored in the (ED). From this model, a planning representation is derived that enables using the model of the environment for navigation purposes.
With use of the ed_navigation plugin [4], an occupancy grid is derived from the world model and published as a nav_msgs/OccupancyGrid. This grid can be used by a motion planner to perform searches in the configuration space of the robot.
With the use of the cb_base_navigation ROS package[5]. The robots are able to deal with end goal constraints. With use of a ROS service, provided by the ed_navigation plugin, an end goal

---

[3] https://github.com/tue-robotics/ed_localization
[4] https://github.com/tue-robotics/ed_navigation
[5] https://github.com/tue-robotics/cb_base_navigation

Fig. 2: A view of the world model created with ED. The figure show the occupation grid as well as (unknown) objects recognized on top of the cabinet.

constraint can be constructed w.r.t. a specific world model entity described by ED. This enables the robot to not only navigate to poses but also to areas or entities in the scene, as illustrated by Figure 3. Somewhat modified versions of the local and global ROS planners available within move_base are used.

(a) Circle constraint, $x^2 + y^2 > a^2$ and $x^2 + y^2 < b^2$, frame *kitchenblock*

(b) Rectangular constraint, $x > a$ and $x < b$ and $y > c$ and $y < d$, frame *livingroom*

Fig. 3: Navigation position constraints w.r.t. other entities in the environment

2.2   Object detection

Detection & Segmentation ED enables integrating sensor data with use of the plugins present in the ed_sensor_integration package. Two different plugins do exist: 1. laser_plugin: Enables tracking of 2D laser clusters. This plugin can be used to track dynamic obstacles such as humans. 2. kinect_plugin: Enables world model updates with use of data from the Microsoft Kinect™. This plugin exposes several ROS services that realize different functionalities:

(a) Segment: Service that segment sensor data that is not associated with other world model entities. Segmentation areas can be specified per entity in the scene. This allows to segment object 'on-top-of' or 'in' a cabinet.

(b) FitModel: Service that fits the specified model in the sensor data of the Microsoft Kinect™. This allows updating semi-static obstacles such as tables and chairs.

The ed_sensor_integration plugins enable updating and creating entities. However, new entities are classified as unknown entities.

(a) Cup, Pills, Diet cok, Fanta - Kinect view   (b) World representation

Fig. 4: ED Perception responsible for the object segmentation and calling the object recognition service. Left, the segmented objects in the robot's sensor frame are displayed; the final annotated world representation is shown at the right picture.

2.3 Object grasping, moving and placing

As for manipulating objects, the architecture is only focused on grasping. The input is the specific target entity in the world model ED. The output is the grasp motion, i.e. joint positions for all joints in the kinematic chain over time. Figure 5 shows the grasping pipeline. A python
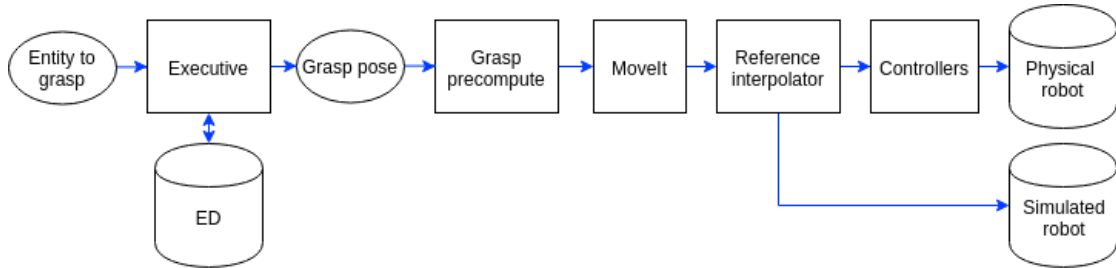

Fig. 5: Custom grasping pipeline base on ED, MoveIt and a separate grasp point determination and approach vector node.

executive queries the current pose of the entity from ED. The resulting grasp pose goes to the grasp precompute component which makes sure that we approach the object in a proper way. MoveIt will produce joint trajectories over time with use of the current configuration, the URDF

model and the final configuration. Note that MoveIt currently does not take any information from ED into account. Finally, the trajectories are sent to the reference interpolator which sends the trajectories either to the controllers or the simulated robot.

The grasping pipeline is extended with an empty spot designator and grasping point determination. The empty spot designator search in an area for an empty spot to place an object by using the occupied area by other objects in this area.

The grasp point determination uses the information about the position and shape of the object in ED to determine the best grasping point. The grasping point is a vector relative to the robot. An example of the determined grasping point is shown in Figure 6.


Fig. 6: Grasping point determination of cylindric object.

3   Image Recognition

In order the classify or train unknown entities, the ed_perception plugin[6] exposes ROS Services to classify the entities in the world model. The ed_perception module interfaces with various image_recognition nodes that apply state of the art image classification techniques based on Convolution Neural Networks (CNN) illustrated in Figure 7.

[6] https://github.com/tue-robotics/ed_perception

Fig. 7: Illustration Convolution Neural Networks (CNN) used in our object recognition nodes with use of Tensorflow.

3.1 Object recognition using Deep Learning

Object recognition is done using Tensorflow: retraining the top-layer of a Inception V3 neural network. The top layers are retrained on a custom dataset using a soft-max top-layer that maps the image representation on a specified set of labels.
In order to create a new training set for specific objects, the ed_perception and the image_recognition packages contains several tools for segmenting and annotating objects. Also tools for retraining neural networks are included.

3.2 Face recognition

Face detection and recognition is done using Openface based on Torch. Openface is an existing state-of-the-art face recognition library. We implemented a ROS node that enables the use of these advanced technologies within the ROS network.

3.3 ROS packages

Our image recognition ROS packages can be found at GitHub[7] with tutorials and documentation. Recently, they have also been added to the ROS Kinetic package list and can be installed as Debian packages:

```
ros−kinetic−image−recognition
```

4 Human-Robot Interface

In order to interact with the robot aside of speech, a web-based Graphical User Interface (GUI) has been designed. The interface has been made with HTML5 and is hosted on the robot itself. This allows multiple users on different platforms (*e.g.* Android, iOS) to access functionality of the robot. The interface is implemented in JavaScript with AngularJS and it offers a graphical interface to the Robot API[8] which exposes all the functionality of the robot. Figure 8 gives an overview of the connections between these components. Figure 9 gives an example of various user

---

[7] https://github.com/tue-robotics/image_recognition
[8] https://github.com/tue-robotics/robot-api

Fig. 8: Overview of the WebGUI architecture. The robot's functionalities are exposed with the Robot API that is implemented in JavaScript. A webserver that is hosting the GUI connects this Robot API to a graphical interface that is offered to multiple clients on different platforms.

interactions that are possible with the GUI and the different commands that can be given to the robot while interacting with the virtual scene.

Fig. 9: Illustration of the 3D scene of the WebGUI. Users can interact with use of the menu that appears when long pressing an object in the scene. On the left figure, the user commands the robot to inspect the selected object, which is the 'cabinet'. When the robot has inspected the 'cabinet', it has found entities on top of it. In the middle figure a grasp command is given to the robot to pick up an object from the cabinet. The last figure show the robot executing that action.

5   Robot Descriptions

5.1   Robot Hardware Descriptions

At the moment of writing of this team description paper, we are not in the possession of the Toyota HSR yet. Therefore we are not able to inform you about any possible additional computing devices in this paper.

5.2   Robot Software Description

An overview of the software used by the Tech United Eindhoven @Home robots can be found in Table 1. All our software is developed open-source at GitHub[9].

Currently, we have some *image_recognition* packages released into the current ROS Kinetic distribution and can be installed with use of *apt*.

Table 1: Software overview of the robots.

Operating system	Ubuntu 16.04 LTS Server
Middleware	ROS Kinetic [1]
Low-level control software	Orocos Real-Time Toolkit [2]
	`https://github.com/tue-robotics/rtt_control_components`
Simulation	Custom kinematics + sensor simulator
	`https://github.com/tue-robotics/fast_simulator`
World model	Environment Descriptor (ED), custom
	`https://github.com/tue-robotics/ed`
Localization	Monte Carlo [3] using Environment Descriptor (ED), custom
	`https://github.com/tue-robotics/ed_localization`
SLAM	Gmapping package
	`http://wiki.ros.org/gmapping`
Navigation	CB Base navigation
	`https://github.com/tue-robotics/cb_base_navigation`
	Global: custom A* planner
	Local: modified ROS DWA [4]
Arm navigation	Custom implementation using MoveIt and Orocos KDL
	`https://github.com/tue-robotics/tue_manipulation`
Object recognition	Tensorflow ROS
	`https://github.com/tue-robotics/image_recognition/tree/`
	`master/tensorflow_ros`
People detection	Custom implementation using contour matching
	`https://github.com/tue-robotics/ed_perception`
Face detection & recognition	Openface ROS
	`https://github.com/tue-robotics/image_recognition/tree/`
	`master/openface_ros`
Speech recognition	Julius Speech Recognition
	`https://github.com/julius-speech/julius`
Speech synthesis	Toyota™ Text-to-Speech
Task executors	SMACH
	`https://github.com/tue-robotics/tue_robocup`

This our current software implementation on our robots AMIGO and SERGIO, which participate in the open league. Because of the pending delivery of the Toyota HSR and related documentation, we are not able to provide the specific software implementation. As described in our selection qualification paper, our goal is to use the same software as possible on all our robots, including the Toyota HSR.

---

[9] `https://github.com/tue-robotics`

5.3   Re-usability of the system for other research groups

Tech United takes great pride in creating and maintaining open-source software and hardware to accelerate innovation. Tech United initiated the Robotic Open Platform website, to share hardware designs. All packages are equipped with documentation and tutorials. Tech United and its scientific staff have the capacity to co-develop (+10 people), maintain and assist with questions.

5.4   Community Outreach and Media

The Tech United team carries out many promotional activities to promote technology and innovation with children. These activities are carried out by separate teams of student assistants. Tech United often visits primary and secondary schools, public events, trade fairs and have regular TV performances. In 2015 and 2016 together, 100+ demos were given and an estimated 50k were reached through live interaction. Tech United has also got a very active (website, and interacts on many social media mediums: Facebook, YouTube, Twitter and Flickr. Our robotics videos are often shared on the IEEE video Friday website.

References

1. Morgan Quigley, Ken Conley, Brian P. Gerkey, Josh Faust, Tully Foote, Jeremy Leibs, Rob Wheeler, and Andrew Y. Ng. ROS: an open-source robot operating system. In *ICRA Workshop on Open Source Software*, 2009.
2. H. Bruyninckx. Open robot control software: the orocos project. In *Proceedings of the 2001 IEEE International Conference on Robotics & Automation*, 2001.
3. D. Fox. Adapting the sample size in particle filters through kld-sampling. *The International Journal of Robotics Research*, 22(12):985–1003, 2003.
4. D. Fox, W. Burgard, and S. Thrun. The dynamic window approach to collision avoidance. *IEEE Magazine on Robotics & Automation*, 4(1):23–33, 1997.